

# Advanced Neural Networks

Applications-Focused Introduction for Beginners  
CMSC 173 - Machine Learning

---

Noel Jeffrey Pinton

October 16, 2025

Department of Computer Science  
University of the Philippines - Cebu

Introduction: Why Advanced Neural Networks?

Convolutional Neural Networks (CNNs)

Generative Models: Creating New Content

GANs: Generative Adversarial Networks

VAEs: Variational Autoencoders

Transformers: The Revolution in AI

Diffusion Models: The Newest Revolution

Ethics & Responsible AI

Key Concept: Attention Mechanism

Using These Models: Practical Guide

Summary & Looking Forward

## Introduction: Why Advanced Neural Networks?

---

# What are Advanced Neural Networks?

## Basic Neural Networks

- Fully connected layers
- Good for tabular data
- Limited to simple patterns
- We learned these already!

## Advanced Architectures

- **CNNs:** For images and spatial data
- **Transformers:** For text and sequences
- **GANs:** Generate new data
- **VAEs:** Learn compressed representations
- **Diffusion:** Create high-quality images

## Why Learn These?

**They power the AI you use every day:**

- ChatGPT (Transformer)
- DALL-E 2 (Diffusion)
- Face unlock on phones (CNN)
- Google Translate (Transformer)
- AI art generators (GAN/Diffusion)

## This Module's Focus

**Understanding applications** rather than complex math!

# Real-World Applications Overview

## CNN Applications: Where You See Them Every Day! ☑

### ☑ Photo Recognition

*Your phone:  
'This is a beach!'*

### ☑ Self-Driving Cars

*Detecting pedestrians,  
stop signs, other cars*

### ☑ Medical Diagnosis

*Finding tumors in  
X-rays and MRI scans*

### ☑ Face Unlock

*Your phone recognizes  
YOUR face to unlock*

### ☑ Image Filters

*Snapchat/Instagram  
fun face filters*

### ☑ Visual Search

*Google Lens:  
Find similar products*

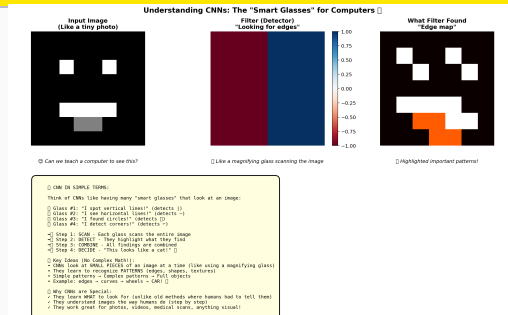
## Course Philosophy

**Learn by seeing what's possible!** We'll focus on understanding what these networks can do and how to use them, not deriving complex mathematics.

## Convolutional Neural Networks (CNNs)

---

# CNNs: What Are They?



## Simple Explanation

CNNs are neural networks designed for images.

They work by:

- Looking at small patches of the image
- Finding patterns (edges, shapes, textures)
- Building up to complex objects
- Making decisions based on what they see

## Key Insight

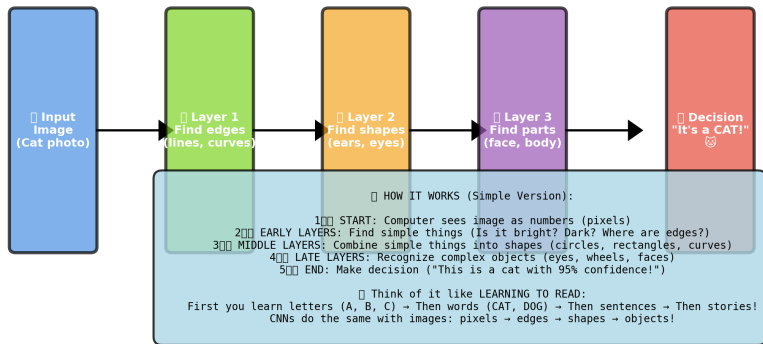
CNNs learn to recognize patterns automatically - no manual feature engineering!

## Why Not Regular NNs?

- Images have too many pixels
- Spatial relationships matter

# How CNNs Process Images

## CNN Architecture: From Simple to Complex (Layer by Layer)



### Processing Pipeline

Input Image → Find Edges → Find Shapes → Find Objects → Decision

### Analogy

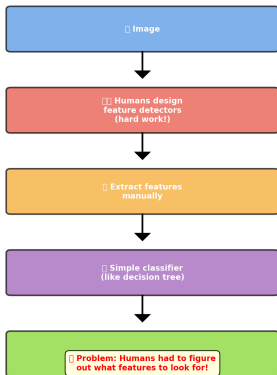
Like how humans see: First we see lines and edges, then shapes, then we recognize "this is a cat!"



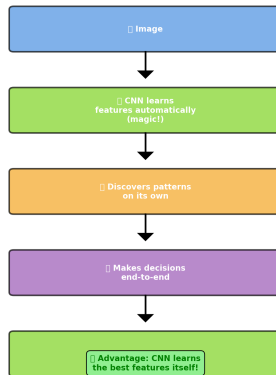
# CNNs vs Traditional Computer Vision

## Why CNNs Changed Everything

### OLD WAY (Before CNNs)



### NEW WAY (CNNs)



## Traditional Methods

- Manual feature design
- Hard to adapt to new tasks
- Limited accuracy
- Lots of expert knowledge needed

## CNNs

- Automatic feature learning
- Easily adapt to new problems
- State-of-the-art accuracy
- Just need training data

## Cancer Detection

### Real Application:

- Detect tumors in X-rays and MRIs
- Classify skin lesions (benign/malignant)
- Analyze mammograms for breast cancer
- Help radiologists work faster

### Impact

- Earlier disease detection
- Fewer missed diagnoses
- Reduced radiologist workload
- Available in rural areas

## Retinal Disease Diagnosis

### Example: Google's Diabetic Retinopathy Detection

- Analyzes eye scans
- Detects diabetes complications
- Matches expert doctor accuracy
- Used in India, Thailand

### Success Story

FDA-approved AI systems now assist doctors in real hospitals!

## Lane Detection

### What CNNs Do:

- Identify road lane markings
- Track lane boundaries in real-time
- Work in various lighting conditions
- Handle curves and intersections

## Object Detection

- Detect pedestrians, cars, cyclists
- Recognize traffic signs and lights
- Estimate distance to objects
- Predict object movement

## Companies Using This

- **Tesla:** Full Self-Driving (FSD)
- **Waymo:** Autonomous taxis
- **Cruise:** Robotaxis in SF
- **Mobileye:** Driver assistance

## Real Deployment

Over 1 million vehicles use CNN-based vision systems today!

## Phone Unlock (Face ID)

### How It Works:

- CNN extracts facial features
- Creates unique "face print"
- Compares to stored template
- Works in different lighting
- Adapts to appearance changes

### Daily Use Cases

- iPhone/Android face unlock
- Photo organization (Google Photos)
- Security access control
- Airport immigration

## Social Media Applications

- **Facebook:** Auto-tag friends in photos
- **Snapchat:** Face filters and effects
- **Instagram:** Beauty filters
- **TikTok:** Face tracking for AR

### Privacy Note

Face recognition raises important privacy concerns - always consider ethics!

## Smart Security Cameras

### Capabilities:

- Detect people vs animals
- Recognize package delivery
- Identify suspicious behavior
- Track movement patterns
- Send targeted alerts

## Consumer Products

- Ring Doorbell cameras
- Nest security systems
- Arlo smart cameras
- Reduce false alarms by 90%

## Retail Applications

### Amazon Go Stores:

- Track what customers pick up
- Automatic checkout (no cashiers)
- Prevent shoplifting
- Analyze shopping behavior

## Industry Impact

Checkout-free stores save 75% of labor costs while improving customer experience!

## Environmental Monitoring

### Applications:

- Track deforestation in Amazon
- Monitor crop health
- Detect illegal fishing
- Assess disaster damage
- Map urban growth

## Real Projects

- **Planet Labs:** Daily Earth imaging
- **Global Fishing Watch:** Ocean monitoring
- **NASA:** Climate change tracking

## Humanitarian Uses

- Count refugees in camps
- Assess natural disaster impact
- Map poverty indicators
- Monitor conflict zones
- Guide relief efforts

## Scale

CNNs can analyze millions of satellite images - impossible for humans alone!

## **Generative Models: Creating New Content**

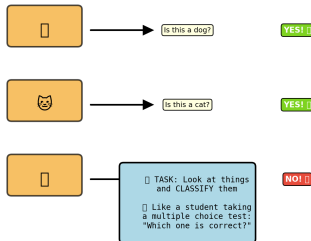
---

# What Are Generative Models?

## Generative vs Discriminative: Creating vs Classifying

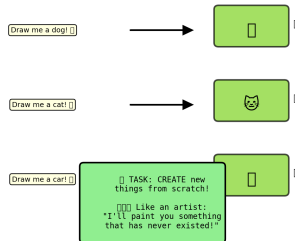
### DISCRIMINATIVE MODELS

(The Classifier - "What is this?")



### GENERATIVE MODELS

(The Creator - "Make me something new!")



## Discriminative Models

### What they do:

- Classify/label existing data
- "Is this a cat or dog?"
- CNNs for image classification

## Generative Models

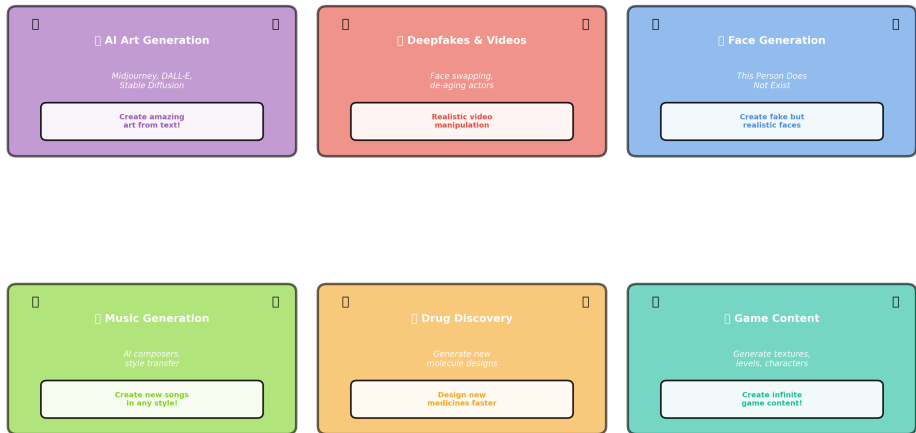
### What they do:

- Create new data
- "Generate a new cat image"
- GANs, VAEs, Diffusion models



# Generative Model Applications Overview

## Generative Models: Creating the Future! 🚀



## Three Main Types We'll Cover

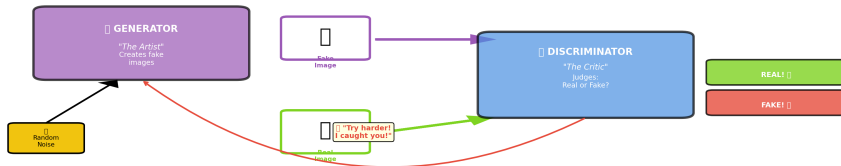
1. **GANs (Generative Adversarial Networks):** Two networks compete to create realistic images

## **GANs: Generative Adversarial Networks**

---

# GANs: The Basic Idea

## GANs: The Artist vs The Critic Game ☐×☐



☐ HOW THE "GAME" WORKS:

- ☐ THE GENERATOR (Artist):
  - Starts with random noise (like random brushstrokes)
  - Creates fake images trying to look real
  - Goal: "Fool the critic into thinking my art is real!"
- ☐ THE DISCRIMINATOR (Critic):
  - Looks at both real images and fake images
  - Tries to tell them apart
  - Goal: "Spot the fakes! Don't be fooled!"
- ☐ THE COMPETITION:
  - 1☐☐ Generator creates a fake image
  - 2☐☐ Discriminator judges: "Real or Fake?"
  - 3☐☐ If caught: Generator learns from mistakes and improves
  - 4☐☐ Discriminator also gets better at detecting fakes
  - 5☐☐ They keep competing until Generator is 50 GOOD, even the expert Discriminator can't tell!
- ☐ THE RESULT:

After thousands of rounds, the Generator becomes an AMAZING artist!  
It can create super realistic images that never existed before!
- ☐ REAL-WORLD ANALOGY:

Think of it like a student artist (Generator) trying to forge a painting,  
while an art expert (Discriminator) tries to detect forgeries.  
The student keeps improving until they're as good as a master!
- ☐ Why GANs are MAGICAL:
  - They learn to create realistic images WITHOUT being told exactly how
  - The competition makes both networks better and better
  - No human has to teach them what "realistic" means - they figure it out!

## Simple Explanation

Two neural networks compete:

- **Generator:** Creates fake images (like an art forger)

## Artbreeder

### What it does:

- Generate unique portraits
- Mix different faces together
- Adjust age, gender, ethnicity
- Create landscapes, album covers
- Used by 10+ million users

### How Artists Use It

- Book cover illustrations
- Character design for games
- Concept art for films
- Social media content

## ThisPersonDoesNotExist.com

- Generates random faces
- 100% synthetic people
- Photorealistic quality
- New face every refresh
- Built with StyleGAN

### Try It Yourself!

Visit the website - every face you see was created by AI, not a photo!

## The Problem

### Malicious Uses:

- Fake celebrity videos
- Misinformation campaigns
- Identity fraud
- Non-consensual content

## The Solution

### GANs fight GANs:

- Train detectors on fake data
- Identify artifacts and inconsistencies
- Real-time video verification
- Protect public figures

## Real Deployments

- **Facebook/Meta:** Deepfake detection system
- **Microsoft:** Video Authenticator tool
- **Intel:** FakeCatcher (96% accuracy)
- **Adobe:** Content Authenticity Initiative

## Arms Race

Detection technology must constantly evolve as GANs improve!

## Why Generate Medical Data?

### Privacy & Scarcity Issues:

- Real patient data is private (HIPAA)
- Rare diseases lack training samples
- Hard to share data between hospitals
- Need diverse examples for AI training

## What GANs Generate

- Synthetic X-rays
- Artificial MRI scans
- Fake patient records
- Privacy-preserving datasets

## Real Research Applications

- **Mayo Clinic:** Generate rare tumor samples
- **Stanford:** Synthetic chest X-rays
- **MIT:** Privacy-safe medical records
- Train better AI without compromising privacy

## Impact

Enables medical AI research while protecting patient privacy!

## Modern Game Development

### How GANs Help:

- Generate unique NPC faces
- Create diverse character variations
- Design textures and materials
- Procedural content generation
- Speed up asset creation

## Real Game Studios

- **EA Sports:** Generate realistic player faces
- **Ubisoft:** NPC diversity in Assassin's Creed
- Reduce manual art time by 70%

## Player Customization

- Infinite character appearance options
- Realistic face generation
- Upload photo for custom avatar
- AI-assisted character design

## Industry Adoption

Major game engines (Unity, Unreal) now integrate GAN-based tools!

## AI Fashion Designers

### What They Generate:

- New clothing designs
- Pattern and texture variations
- Color scheme combinations
- Style transfer between eras
- Personalized recommendations

## Fashion Companies Using AI

- **Stitch Fix:** Personalized designs
- **Tommy Hilfiger:** IBM collaboration
- **Zalando:** Generated fashion models

## Virtual Try-On

- Generate how clothes look on you
- Try outfits without physically wearing
- Reduce online shopping returns
- Personalized styling suggestions

## Business Impact

AI-designed collections sell out 30% faster than traditional designs!




## **VAEs: Variational Autoencoders**

---

# VAEs: What Are They?

## VAE: Learning to Compress and Create

### VAE: The Compression Artist



VAE IN SIMPLE TERMS (Variational Autoencoder):

THINK OF IT LIKE ZIP FILES:  
When you zip a large file, it gets compressed into something smaller.  
VAEs do the same thing with images, but in a SMART way!

HOW IT WORKS:

1. ENCODER (The Compressor):
  - Takes a big image (like a photo of a cat)
  - Squeezes it into a tiny "secret code" (just a few numbers!)
  - Like compressing a photo to save space on your phone
2. LATENT SPACE (The Secret Code):
  - The compressed representation
  - Contains the ESSENCE of the image (not every pixel, just the important info)
  - Like a recipe card instead of the full meal
3. DECODER (The Decompressor):
  - Takes the secret code
  - Recreates the image from memory
  - Like unzipping a file, but it GENERATES the image

THE MAGIC TRICK:  
Once trained, you can:

- Create NEW images by giving it NEW random codes!
- Mix codes to blend images (cat + dog = ???)
- Change codes slightly to modify images

REAL-WORLD ANALOGY:  
Imagine describing a face with just a few words: "round face, blue eyes, blonde hair, smiling"  
From just that SHORT description, someone could draw a complete face!  
That's what VAE's secret code does!

VAE vs GAN:

- VAE: More stable, creates smooth variations, great for compression
- GAN: More realistic, but training is trickier (needs the competition)

USES:

- Generating new faces, artwork, designs
- Image compression (better than JPEG!)
- Anomaly detection (if compression fails, something's weird!)
- Creating variations of existing images

## Simple Explanation

VAEs compress data into a small code, then decompress it:

## Manufacturing Quality Control

### How It Works:

- Train VAE on normal products
- VAE learns what "normal" looks like
- Defects reconstruct poorly
- High reconstruction error = defect!

### Real Applications

- Detect scratches on surfaces
- Find cracks in materials
- Identify missing components
- Automated quality inspection

## Other Anomaly Detection Uses

- **Cybersecurity:** Detect network intrusions
- **Finance:** Identify fraudulent transactions
- **Healthcare:** Flag unusual patient vitals
- **IoT:** Detect sensor failures

### Advantage

Works without labeled defect examples - learns from normal data only!

## Why VAEs for Compression?

### Advantages over JPEG:

- Better quality at low bitrates
- Learned compression (adapts to content)
- Can compress to tiny sizes
- Semantic preservation

### How It Works

- Encoder compresses to latent code
- Store only the small code
- Decoder reconstructs when needed
- 10-100x smaller than JPEG

## Real-World Uses

- Store medical imaging archives
- Stream video at lower bandwidth
- Compress satellite imagery
- Mobile app image caching

## Research Example

Google's neural image compression beats JPEG by 50% in quality metrics!

## Pharmaceutical Discovery

### Traditional Approach:

- Test millions of molecules
- Takes 10+ years per drug
- Costs billions of dollars
- High failure rate

### VAE Approach

- Learn from existing drugs
- Generate similar molecules
- Optimize for target properties
- Find candidates much faster

## Real Pharmaceutical AI

- **Insilico Medicine:** Generated novel molecules
- **Atomwise:** AI drug discovery platform
- **BenevolentAI:** COVID-19 drug repurposing
- Reduce discovery time by 75%

## Major Milestone

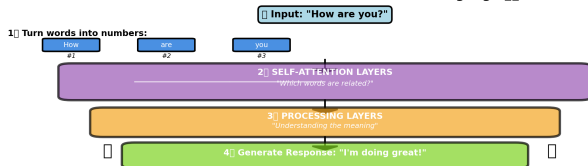
First AI-discovered drug entered human trials in 2020!

## **Transformers: The Revolution in AI**

---

# Transformers: What Are They?

## Transformers: How ChatGPT Understands Language



### TRANSFORMER IN SIMPLE TERMS:

- OLD APPROACH (RNNs): Read sentence word-by-word, left to right
  - Like reading a book one word at a time
  - Slow and forgets earlier words

• Example: By the time you read "dog", you forgot the "big brown" at the start

### TRANSFORMER APPROACH: Look at ALL words at once!

- Like looking at the whole sentence together
- Uses ATTENTION to see which words are related
- Fast and remembers everything

### THE MAGIC - SELF-ATTENTION:

When processing the word "it", the transformer looks at the whole sentence:  
"The cat sat on the mat because it was tired"  
→ Attention figures out "it" refers to "cat", not "mat"!

### HOW CHATGPT USES THIS:

- You type: "Translate 'hello' to Spanish"
- Transformer reads ALL words simultaneously
- Attention connects "translate" with "hello" and "Spanish"
- It understands: [ACTION: translate] [FROM: hello] [TO: Spanish]
- Generates: "Hola"

### WHY TRANSFORMERS CHANGED AI:

- Can be trained on MASSIVE amounts of text (entire internet!)
  - Understands context better (knows "bank" = money vs river)
- Much faster than old methods (processes sentences in parallel)
- Can handle very long text (remember things from paragraphs ago)

### KEY INSIGHT:

Transformers are like having a super-smart friend who can:

- Read an entire book in seconds
- Remember every detail
- Understand how everything connects
- Explain it back to you in simple terms

## Simple Explanation

Transformers process sequences by paying attention to relevant parts:

- Designed for text, but work on images/audio too

# Transformer Applications Overview

## Transformers in Daily Life: You Use These Every Day! 📱



## Why Transformers Changed Everything

Before 2017: RNNs struggled with long sequences. After 2017: Transformers enabled GPT, BERT, and the current AI revolution!



## What ChatGPT Can Do

### Capabilities:

- Answer questions
- Write code and debug
- Compose essays and emails
- Explain complex topics
- Translate languages
- Creative writing

## Real Usage Statistics

- 100+ million weekly users
- Fastest-growing consumer app
- Used in 185+ countries

## How Students Use It

- Homework help and tutoring
- Research assistance
- Programming debugging
- Study guide creation
- Language learning
- Career advice

## Built With Transformers

GPT-4 uses a massive transformer with 175+ billion parameters!

## Old vs New Approach

### Before Transformers (2016):

- Phrase-based translation
- Limited context understanding
- Often awkward output

### After Transformers (2017+):

- Sentence-level context
- Natural, fluent translations
- 60% reduction in errors

## Features Powered by Transformers

- 133 languages supported
- Real-time conversation mode
- Camera translation (point and translate)
- Offline translation
- Context-aware results

## Daily Impact

500+ million people use Google Translate every day!

## AI Pair Programmer

### What Copilot Does:

- Suggests code as you type
- Writes entire functions
- Explains existing code
- Converts comments to code
- Generates tests
- Fixes bugs

## Real Developer Impact

- 46% of code written by AI
- 55% faster task completion
- Used by 1.2 million developers

## How It Works

- Built on GPT (Codex model)
- Trained on billions of lines of code
- Understands context from your files
- Suggests in real-time
- Supports 12+ programming languages

## For Students

Great learning tool - see how experts solve problems!

## Gmail Smart Compose

### Features:

- Suggests next words/sentences
- Learns your writing style
- Adapts to context
- Multi-language support
- Works on mobile too

### Time Savings

- Average user saves 1 billion characters/week
- Reduces writing time by 11%
- 4+ billion emails use it daily

## Other Email AI Features

- **Smart Reply:** Suggest full responses
- **Subject suggestions:** Auto-generate subjects
- **Tone adjustment:** Make emails more formal
- **Grammar correction:** Fix mistakes

## All Powered by Transformers

These "small" conveniences use the same tech as ChatGPT!

## Automatic Summarization

### What It Does:

- Read long documents
- Extract key points
- Generate concise summary
- Preserve important details
- Save reading time

## Real Products

- **Microsoft Word:** Auto-summarize
- **Slack:** Thread summaries
- **Notion AI:** Note summarization
- **Chrome extensions:** Web page summaries

## Use Cases

- Research paper summaries
- News article digests
- Legal document review
- Meeting notes condensation
- Customer feedback analysis


## Productivity Boost

Lawyers using AI summarization save 60% of document review time!

# Vision Transformers: Images Meet Transformers

## Vision Transformers: Teaching Language Models to See! 🐱

🐱 THE BIG IDEA: Treat images like sentences!



### 🐱 VISION TRANSFORMERS (ViT) IN SIMPLE TERMS:

#### 🐱 THE CLEVER TRICK:

Instead of inventing a NEW architecture for images, we realized:  
"Why not use the SAME transformer that works great for text?"

#### 🐱 HOW IT WORKS:

- 1🐱 Cut image into small squares (patches) - like cutting a pizza into slices
- 2🐱 Flatten each patch into a list of numbers (treat it like a "word")
- 3🐱 Feed these "image words" into a regular transformer
- 4🐱 Let attention figure out which patches are related!

#### 🐱 EXAMPLE:

If you have a photo of a cat:

- Patch 1: Contains part of the ear
- Patch 2: Contains the eye
- Patch 3: Contains whiskers

→ Attention learns: "These patches together = cat face!"

#### 🐱 ViT vs CNN:

- CNN: Looks at local neighborhoods (like looking through a magnifying glass)
- ViT: Looks at the whole image at once (like seeing the big picture)
- ViT: Can find connections between far-apart parts of the image!

#### 🐱 WHY IT'S AWESOME:

- 🐱 Simpler architecture than CNNs (one model for everything!)
- 🐱 Can handle huge images (just use more patches)
- 🐱 Learns better from LOTS of data
- 🐱 Same code works for images, text, video, audio!

#### 🐱 THE FUTURE:

- Multi-modal models like GPT-4, Claude use this idea:
  - They can understand BOTH text AND images together!
  - Ask: "What's in this photo?" and show a picture
  - AI can answer because it treats images and text the same way!

## Vision Transformers (ViT)

### Applying transformers to images:

- Break image into patches (like words)

## **Diffusion Models: The Newest Revolution**

---


# Diffusion Models: How They Work

## Diffusion Models: The Magic of "Denoising" ☐☐

→ FORWARD PROCESS: Gradually Add Noise (Training)



← REVERSE PROCESS: Gradually Remove Noise (Generation)



### ☐ DIFFUSION MODELS IN SIMPLE TERMS:

☐ THE CORE IDEA: Learn to reverse the process of adding noise!

☐ ANALOGY - Like a Magic Eraser:

Imagine you have a beautiful drawing, but someone scribbled all over it with random marks.  
A diffusion model is like a magic eraser that can gradually remove the scribbles until the original drawing appears again!

☐ TWO PHASES:

1☐☐ TRAINING (Forward): Add noise step by step

- Start with real images (like photos of cats)
- Gradually add random noise until it's completely fuzzy
- Save each step to learn the pattern

2☐☐ GENERATION (Reverse): Remove noise step by step

- Start with pure random noise (like TV static)
- AI predicts: "What does this look like with LESS noise?"
- Remove a little noise at a time
- Eventually reveals a clear image!

☐ HOW THE AI LEARNS:

During training, the AI learns: "Given a noisy image, what's the clean version?"  
It practices this over and over with millions of images.  
Then during generation, it applies this skill repeatedly to create new images!

☐ WHY IT'S SPECIAL:


- Very stable training (unlike GANs that can be tricky)
- High quality results (super realistic images!)
- Controllable (you can guide what it creates with text)
- Flexible (works for images, audio, video, 3D models)

☐ THE MAGIC:

You start with random noise (complete chaos) and step-by-step, the AI sculpts it into a beautiful image - like a sculptor revealing a statue hidden inside a block of marble!



## Generative Models Comparison: Which One to Use? 🤖



### DETAILED COMPARISON:

Feature	GANs	VAEs	Diffusion
Image Quality			
Training Stability			
Generation Speed			
Diversity			
Control (text)			

**WHEN TO USE EACH:**

- GANs:** When you need fast generation and have lots of training data  
Example: Generating realistic faces, style transfer
- VAEs:** When you need smooth interpolation and stable training  
Example: Data compression, anomaly detection, simple generation
- DIFFUSION:** When you need the highest quality and text control  
Example: AI art (DALL-E, Midjourney), image editing, creative applications

**CURRENT WINNER: Diffusion Models!**  
They're the technology behind most modern AI art tools because they:

- Generate the highest quality images
- Work great with text prompts
- Are stable and reliable to train
- Can be controlled and edited easily

**THE TREND:**  
Diffusion models have largely replaced GANs for image generation!  
They're slower, but the quality difference is worth it for most applications.

### GANs

**Pros:** Fast generation

### VAEs

**Pros:** Stable, good latent space

### Diffusion

**Pros:** Best quality, stable

# Diffusion Applications Overview

## Diffusion Models: Creating the Impossible! 🪄



## Why Diffusion Models Won

They power DALL-E 2, Midjourney, Stable Diffusion - the best AI image generators today!

## What DALL-E 2 Can Do

### Text-to-Image Generation:

- Type a description, get an image
- Photorealistic or artistic styles
- Combine multiple concepts
- Edit existing images
- Outpainting (extend images)

### Example Prompts

- "A cat astronaut on Mars"
- "Oil painting of a sunset over Manila"
- "Teddy bear shopping for groceries"

## Real-World Uses

- Marketing content creation
- Concept art for entertainment
- Educational illustrations
- Social media graphics
- Product mockups

## By OpenAI

Same company behind ChatGPT - 1.5+ million users create images daily!

## What Makes Midjourney Special

### Artistic Focus:

- Exceptionally beautiful outputs
- Strong artistic style
- Great for fantasy/sci-fi art
- Discord-based interface
- Community of 16+ million users

### Popular Use Cases

- Book cover designs
- Album artwork
- Game concept art
- NFT art generation

## Industry Impact

- Artists use it for inspiration
- Magazine covers created with AI
- Award-winning art competitions
- Commercial illustration work

### Controversy

AI art won Colorado State Fair - sparked debate about AI creativity!

## Why Stable Diffusion is Different

### Open Source:

- Free to use and modify
- Run on your own computer
- Customize and fine-tune
- No usage restrictions
- Active developer community

## Technical Details

- Can run on consumer GPUs
- Faster than DALL-E 2
- Extensible with plugins
- Multiple versions and variants

## Popular Applications Built With It

- DreamStudio (official interface)
- Automatic1111 (popular UI)
- ComfyUI (node-based editor)
- Mobile apps (Draw Things)
- Photoshop plugins

## Democratizing AI

Anyone with a decent computer can now generate professional-quality images!

## Professional Image Editing

### Firefly Features:

- Text-to-image generation
- Generative fill (edit parts of images)
- Text effects (3D text styles)
- Generative recolor
- Integrated in Photoshop

### Key Advantages

- Trained on Adobe Stock (licensed data)
- Commercially safe to use
- Professional quality outputs
- Seamless Creative Cloud integration

## Real Designer Workflows

- Remove unwanted objects
- Extend backgrounds
- Generate variations quickly
- Create mockups from descriptions
- Speed up creative process 10x

## Industry Standard

Adobe's AI tools are becoming essential for professional designers!

## Text-to-Video AI

### Emerging Applications:

- Generate short video clips
- Animate static images
- Create transitions
- Style transfer for video
- AI-assisted editing

### Current Platforms

- **Runway Gen-2:** Text-to-video
- **Pika Labs:** Video generation
- **Stable Video Diffusion:** Open source

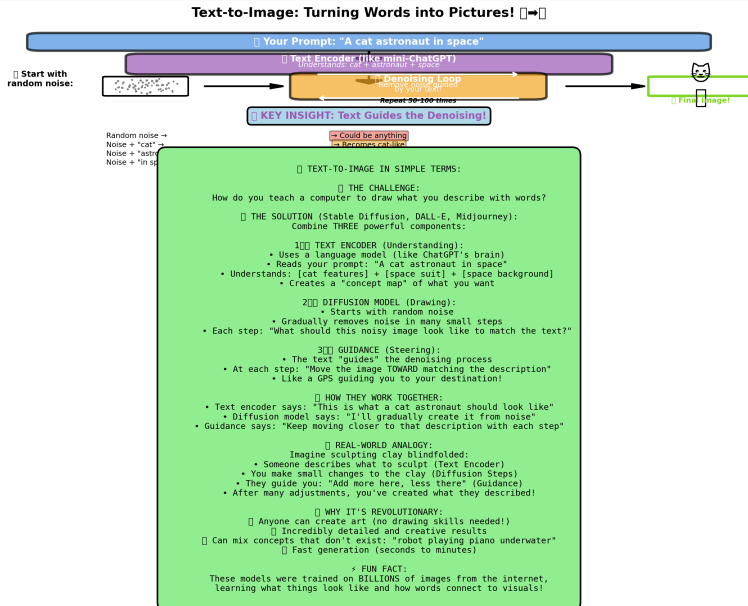
## Use Cases

- Social media content
- Marketing videos
- Animated presentations
- Film pre-visualization
- Game cinematics

## Future is Coming

Video generation is improving rapidly - expect major breakthroughs soon!

# Text-to-Image Process Explained



## How It All Works Together

### 1. Text Encoder (Transformer): Understands your description



## **Ethics & Responsible AI**

---

## 📖 Ethics: The Responsible Use of AI Generation

### POSITIVE USES

□ Creative Expression	Artists using AI as a tool
□ Medical Research	Generating synthetic medical data
□ Education	Creating learning materials
& Accessibility	Helping people with disabilities
□ Scientific Discovery	Simulating molecules, <del>scenarios</del>

### CONCERNS & RISKS

□ Deepfakes	Fake videos spreading misinformation
□ Copyright Issues	Who owns AI-generated art?
□ Job Displacement	Will AI replace human artists?
□ Privacy Violations	Generating fake images of real people
□ Fake News	Creating fake but realistic content

#### □ HOW TO USE GENERATIVE AI RESPONSIBLY:

- 1□ BE TRANSPARENT: Always disclose when content is AI-generated  
Example: Label AI art as "Created with AI" or "AI-assisted"
  - 2□ RESPECT CONSENT: Don't create fake images/videos of real people without permission  
Example: Don't deepfake someone's face onto inappropriate content
  - 3□ VERIFY SOURCES: Don't believe everything you see online  
Example: Check if that viral video might be AI-generated
  - 4□ CREDIT PROPERLY: Acknowledge both the AI tool AND human creativity  
Example: "Concept by [Artist], generated with [AI Tool]"
  - 5□ USE FOR GOOD: Think about the impact of what you create  
Example: Use AI to help people, not to deceive or harm them
- REMEMBER: Just because we CAN create something doesn't mean we SHOULD!
- THE GOAL: Use AI as a tool to enhance human creativity and solve problems, not to deceive, harm, or replace human connection.

## Important Questions to Consider

As these technologies become powerful, we must think carefully about their impact!

## Misinformation & Deepfakes

### Concerns:

- Fake news and propaganda
- Identity fraud
- Non-consensual content
- Erosion of trust in media

### Solutions:

- Detection technology
- Digital watermarking
- Media literacy education
- Legal frameworks

## Bias & Fairness

### Problems:

- Biased training data
- Perpetuating stereotypes
- Unfair representation
- Discrimination in outputs

### Mitigation:

- Diverse training datasets
- Bias testing and auditing
- Responsible AI guidelines
- Inclusive development teams

## Copyright & Intellectual Property

### Questions:

- Who owns AI-generated content?
- Is training on copyrighted data fair use?
- Should artists be compensated?
- How to attribute AI creations?

### Current Debates:

- Ongoing lawsuits (artists vs AI companies)
- New legislation being proposed
- Industry opt-out mechanisms

## Job Displacement

### Concerns:

- Will AI replace creative jobs?
- Impact on artists, writers, designers
- Economic inequality
- Need for reskilling

### Opportunities:

- AI as a tool, not replacement
- New creative possibilities
- Democratization of creation
- Focus on uniquely human skills

## Your Responsibility

As future AI practitioners, think critically about the impact of your work!

## **Key Concept: Attention Mechanism**

---

# Understanding Attention

## Attention Mechanism: Focusing on What Matters


OLD WAY: Reading Every Word Equally



Problem: Computer treats ALL words as equally important! 😞

Question: "What sat on the mat?" → Hard to know "cat" is the answer!

NEW WAY: Attention! Focus on Important Words



Attention lets the model FOCUS on important words!

ATTENTION IN SIMPLE TERMS:

When you read "The cat sat on the mat", you naturally focus more on "cat" and "mat" because they're the important nouns. Words like "the" and "on" are less important.

ATTENTION teaches AI to do the same thing - focus on what matters!

## What is Attention?

A mechanism that lets neural networks focus on relevant parts:

- In text: Focus on important words in a sentence

## Problem Without Attention

Translating: "The cat sat on the mat"

Old approach:

- Process word by word left to right
- Forget earlier context
- Struggle with long sentences
- Poor word alignment

## With Attention

For each output word, the model:

- Looks at ALL input words
- Focuses on relevant ones
- "sat" pays attention to "cat" and "mat"
- Handles long-distance dependencies
- Better translation quality

## Why It's Revolutionary

Attention enabled Transformers to outperform all previous architectures!

## **Using These Models: Practical Guide**

---



## Free/Accessible Tools

Try these today:

- **ChatGPT:** Free tier available
- **Bing Image Creator:** Free DALL-E access
- **Google Colab:** Run Stable Diffusion free
- **Hugging Face:** Try many models online
- **Runway:** Free trial for video

## Learning Resources

- Fast.ai courses (free)
- Hugging Face tutorials
- Papers with Code
- YouTube: Two Minute Papers

## For Developers

Build your own:

- PyTorch or TensorFlow
- Hugging Face Transformers library
- Stable Diffusion on GitHub
- Pre-trained models available
- Fine-tune on your data

## Start Small

Use existing models before building from scratch - learn by doing!

## Writing Good Prompts

### Be specific:

- Describe style (photorealistic, cartoon, oil painting)
- Specify details (colors, lighting, mood)
- Mention composition (close-up, wide shot)
- Add quality keywords (4K, detailed, masterpiece)

### Example Good Prompt

"A majestic golden retriever sitting in a flower meadow at sunset, photorealistic, warm lighting, shallow depth of field, 4K quality"

## Iteration is Key

- Generate multiple variations
- Refine your prompt
- Use negative prompts (what to avoid)
- Adjust parameters (steps, guidance)
- Learn from community prompts

### Pro Tip

Check out prompt libraries (Lexica.art, PromptHero) to learn from others!

## Challenge: Poor Results

### If outputs look bad:

- Improve your prompt specificity
- Try different seed values
- Adjust generation parameters
- Use a different model/variant
- Increase generation steps

## Challenge: Wrong Anatomy/Details

### Known limitations:

- Hands and fingers often wrong
- Text in images unclear
- Physics may be incorrect
- Use inpainting to fix specific parts

## Challenge: Slow Generation

### Speed up:

- Use lower resolution first
- Reduce number of steps
- Try faster samplers
- Use GPU acceleration
- Consider paid services for speed

## Challenge: Reproducibility

### Get consistent results:

- Save your seed numbers
- Keep prompt exactly the same
- Note all parameters used
- Use img2img for variations

## Summary & Looking Forward

---

## What We Learned

### Five major architectures changing the world:

1. **CNNs:** Revolutionized computer vision (medical imaging, self-driving cars, face recognition)
2. **GANs:** Generate realistic images (AI art, deepfakes, synthetic data)
3. **VAEs:** Compress and generate (anomaly detection, drug discovery)
4. **Transformers:** Dominated NLP (ChatGPT, translation, code generation)
5. **Diffusion:** Best image generation (DALL-E 2, Midjourney, Stable Diffusion)

## Main Message

These aren't just research projects - they're tools you can use TODAY in real applications!

## CNNs Applications

- Medical tumor detection
- Self-driving lane detection
- Phone face unlock
- Security cameras
- Satellite imagery analysis

## GAN Applications

- Artbreeder AI art
- Deepfake detection
- Synthetic medical data
- Game character creation
- Fashion design

## Transformer Applications

- ChatGPT conversations
- Google Translate
- GitHub Copilot
- Email auto-complete
- Document summarization

## Diffusion Applications

- DALL-E 2 image generation
- Midjourney art creation
- Stable Diffusion (open source)
- Adobe Firefly editing
- Video generation (emerging)

## Trends to Watch

Next 1-2 years:

- **Multimodal AI:** Text, image, audio, video together
- **Better video generation:** Movie-quality AI videos
- **3D generation:** Create 3D models from text
- **Real-time generation:** Instant results
- **Personalization:** AI that learns your style

## Career Opportunities

Skills in demand:

- AI/ML engineering
- Prompt engineering
- AI safety and ethics
- Creative AI applications
- AI product management

## Get Involved

The best way to learn is to experiment - start building today!

## Hands-On Practice

- Try Stable Diffusion on Colab
- Build projects with Hugging Face
- Fine-tune models on your data
- Participate in Kaggle competitions
- Contribute to open source projects

## Online Courses

- Fast.ai: Practical Deep Learning
- Stanford CS230: Deep Learning
- Coursera: Deep Learning Specialization
- Hugging Face NLP Course (free)

## Stay Updated

- Follow Papers with Code
- Read AI newsletters (The Batch, etc.)
- Watch Two Minute Papers (YouTube)
- Join AI Discord communities
- Attend local meetups

## Next Steps in This Course

**Workshop:** Hands-on coding with ResNet, GPT-2, Stable Diffusion - let's use these models!



**Thank you for your attention!**

### Contact Information

**Instructor:** Noel Jeffrey Pinton

**Course:** CMSC 173 - Machine Learning

**Institution:** University of the Philippines - Cebu

**Department:** Computer Science

### Remember

Advanced neural networks are tools that empower creativity and solve real problems. Use them responsibly and ethically!